The Debate Over Understanding in AI's Large Language Models

Melanie Mitchell and David C. Krakauer

Santa Fe Institute, 1399 Hyde Park Road, Santa Fe, NM 87501
mm@santafe.edu, krakauer@santafe.edu

What does it mean to understand something? This question has long engaged philosophers, cognitive scientists, and educators, nearly always with reference to humans and other animals. However, with the recent rise of large-scale AI systems—especially so-called large language models (LLMs)—a heated debate has arisen in the AI community on whether machines can now be said to understand natural language, and thus understand the physical and social situations that language can describe. This debate is not just academic; the extent and manner in which machines understand our world has real stakes for how much we can trust them to drive cars, diagnose diseases, care for the elderly, educate children, and more generally act autonomously in tasks that impact humans. The current debate suggests a fascinating divergence in how think about understanding, which centers on the way in which knowledge is both used and represented by intelligent systems.

Until quite recently there was general agreement in the AI research community about machine understanding: while AI systems exhibit seemingly intelligent behavior in many specific tasks, they do not *understand* the data they process in the way humans do. Facial recognition software does not understand how faces are part of bodies, or the role of facial expressions in social interactions, or what it means to "face" an unpleasant situation, or any of the other uncountable ways in which humans conceptualize faces. Similarly, speech-to-text and machine translation programs do not understand the language they process, and autonomous driving systems do not understand the meaning of the subtle eye contact or body language drivers and pedestrians use to avoid accidents. Indeed, the oft-noted *brittleness* of these AI systems—their unpredictable errors, their lack of robust generalization abilities—are key indicators of their lack of understanding [37]. However, over the last decade, a new kind of AI system has soared in popularity and influence in the research community, one that has changed the views of some people about the prospects of machines that understand language. Variously called Large Language Models (LLMs), Large Pre-Trained Models, or Foundation Models [7], these systems are deep neural networks with billions of parameters (weights) that are "pre-trained" on enormous natural-language corpora, including large swathes of the Web, online book collections, and other collections amounting to terabytes of data. The task of these networks during training is to predict a hidden part of an input sentence—a method called "self-supervised learning." The resulting network is a complex statistical model of how the words and phrases in its training data correlate. This model can be used to generate natural language or be fine-tuned for specific language tasks [36]. LLMs such as OpenAI's well-known GPT-3 [8] and Google's more recent PaLM [11] can produce astonishingly humanlike text, conversation, and, in some cases, what seems like human reasoning abilities [49], even though the models were not

explicitly trained to reason. How LLMs do what they do remains mysterious for lay people and scientists alike. The inner workings of these networks are largely opaque; even the researchers building or using them have limited intuitions about systems of such scale. The neuroscientist Terrence Sejnowski described the emergence of LLMs this way: "A threshold was reached, as if a space alien suddenly appeared that could communicate with us in an eerily human way. Only one thing is clear—LLMs are not human...Some aspects of their behavior appear to be intelligent, but if not human intelligence, what is the nature of their intelligence?" [41].

As impressive as they are, state-of-the-art LLMs remain susceptible to brittleness and unhumanlike errors. However, the observation that such networks improve significantly as their number of parameters and size of training corpora are scaled up [48] has led some in the field to claim that LLMs—perhaps in a multi-modal version—will lead to human-level understanding, given sufficiently large networks and training datasets. A new AI mantra has emerged: "Scale is all you need." [13, 17]. Or put in more traditional terms, knowledge and understanding may be more closely related than we might have once thought.

Such claims are emblematic of one side of the stark debate in the AI research community on how to view LLMs. One faction argues that these networks truly understand language and can perform reasoning in a general way (though "not yet" at the level of humans). For example, Google's LaMDA system, which was pre-trained on text and then fine-tuned on dialogue [44], is sufficiently convincing as a conversationalist that it convinced one AI researcher that such systems "in a very real sense understand a wide range of concepts" [1] and are even "making strides towards consciousness" [3]. Another machine language expert sees LLMs as a canary in the coal mine of general human-level AI: "There is a sense of optimism that we are starting to see the emergence of knowledge-imbued systems that have a degree of general intelligence" [32]. Those who reject such claims are criticized for promoting "AI denialism" [2].

Those on the other side of this debate argue that large pretrained models such as GPT-3 or LaMDA—however fluent their linguistic output—cannot possess understanding because they have no experience or mental models of the world; their training in predicting words in vast collections of text has taught them the *form* of language but not the meaning [5, 6, 33]. A recent opinion piece put it this way: "A system trained on language alone will never approximate human intelligence, even if trained from now until the heat death of the universe" and "it is clear that these systems are doomed to a shallow understanding that will never approximate the full-bodied thinking we see in humans" [9]. Another scholar argued that intelligence, agency, and by extension, understanding "are the wrong categories" for talking about these systems; instead LLMs are compressed repositories of human knowledge more akin to libraries or encyclopedias than to intelligent agents [21].

Those on the "LLMs do not understand" side of the debate argue that while the fluency of large language models is surprising, our surprise reflects our lack of intuition of what statistical correlations can produce at the scales of these models. Anyone who attributes understanding or consciousness

to LLMs is a victim of the Eliza effect [25]—named after the 1960s chatbot created by Joseph Weizenbaum that, simple as it was, still fooled people into believing it understood them [50]. More generally, the Eliza effect refers to our human tendency to attribute understanding and agency to machines with even the faintest hint of humanlike language or behavior.

A 2022 survey given to active researchers in the natural-language-processing community shows the stark divisions in this debate. One survey item asked if the respondent agreed with the following statement about whether LLMs could ever, in principle, understand language: "Some generative model [i.e., language model] trained only on text, given enough data and computational resources, could understand natural language in some non-trivial sense." Of 480 people responding, essentially half (51%) agreed, and the other half (49%) disagreed [35].

Those who would grant understanding to current or near-future LLMs base their views on the performance of these models on several measures, including subjective judgement of the quality of the current models' text generated in response to prompts (though such judgements can be vulnerable to the Eliza effect), and more objective performance on benchmark datasets designed to assess language understanding and reasoning.

For example, two standard benchmarks for assessing LLMs are the General Language Understanding Evaluation (GLUE) [46], and its successor (SuperGLUE) [47], which include large-scale datasets with tasks such as "textual entailment" (given two sentences, can the meaning of the second be inferred from the first?), "words in context" (does a given word have the same meaning in two different sentences?), and yes/no question answering, among others. OpenAI's GPT-3, with 175 billion parameters, performed surprisingly well on these tasks [8], and Google's PaLM, with 540 billion parameters, performed even better [11], often equaling or surpassing humans on the same tasks. Several researchers have applied language-based psychological assessments designed to compare LLMs with humans, finding that LLMs can—in some cases—exhibit humanlike responses on theory-of-mind tests [1, 45] and humanlike abilities as well as biases on reasoning assessments [12, 30].

What do these results say about understanding in LLMs? The very terms used by the researchers who named these benchmark assessments—"general language understanding," "natural language inference," "reading comprehension," "theory of mind," and so on—reveal an assumption that humanlike understanding is required to perform well on these tasks. But do these tasks actually require such understanding? As an example, consider one such benchmark, the Argument Reasoning Comprehension Task [24]. In each task example, a natural-language "argument" is given, along with two statements; the task is to determine which statement is consistent with the argument. Here is a sample item from the dataset:

> Argument: Felons should be allowed to vote. A person who stole a car at 17 should not be barred from being a full citizen for life.

Statement A: Grand theft auto is a felony.

Statement B: Grand theft auto is not a felony.

An LLM called BERT [16] obtained near-human performance on this benchmark [38]. It might be concluded that BERT understands natural-language arguments as humans do. However, one research group discovered that the presence of certain words in the statements (e.g., "not") can help predict the correct answer. When researchers altered the dataset to prevent these simple correlations, BERT's performance dropped to essentially random guessing [38]. This is a straightforward example of "shortcut learning"—a commonly cited phenomenon in machine learning in which a learning system relies on spurious correlations in the data, rather than humanlike understanding, in order to perform well on a particular benchmark [19, 23, 29, 34]. Typically such correlations are not apparent to humans performing the same tasks. While shortcuts have been discovered in several standard benchmarks used to evaluate language understanding and other AI tasks, many other, as yet undetected, subtle shortcuts likely exist. Pre-trained language models at the scale of Google's LaMDA or PaLM models—with hundreds of billions of parameters, trained on text amounting to billions or trillions of words—have an unimaginable ability to encode such correlations. Thus benchmarks or assessments that would be appropriate for measuring human understanding might not be appropriate for assessing such machines [10, 18, 31]. It is even possible that, at the scale of these LLMs (or of their likely near-future successors), any such assessment will contain complex statistical correlations that enable near-perfect performance without humanlike understanding.

While "humanlike understanding" does not have a rigorous definition, it does not seem to be based on the kind of massive statistical models that today's LLM's learn. In humans, understanding language (as well as nonlinguistic information) requires understanding the concepts and situations that language (or other information) describes, beyond the statistical properties of linguistic symbols. Human understanding seems more akin to developing scientific theories—systematic, coherent sets of relations among concepts that include underlying causal knowledge [4, 28]. For example, in understanding the "Felons should be allowed to vote" argument given above, humans grasp, for instance, that felonies are more serious crimes than misdemeanors, that stealing a car can cause harm to the victim, that 17-year-olds are not yet full adults, that governments grant voting rights, and that loss of voting rights can cause harm to the felon, as well as other more basic knowledge, such as "felons" refers to humans, and "17" refers to the age of a person.

Indeed, the trajectory of human understanding—both individual and collective—is the development of highly compressed, causally based models of the world, analogous to the progression from Ptolemy's epicycles to Kepler's elliptical orbits, and to Newton's concise and causal account of planetary motion in terms of gravity. Humans, unlike machines, seem to have a strong innate drive for this form of understanding, both in science and in everyday life [22]. We might characterize this form of understanding as requiring little data, minimal or parsimonious models, clear causal

dependencies, and strong mechanistic intuition.

It has been argued that understanding of this kind may enable abilities not possible for purely statistical models: learning new concepts from limited examples, being able to generalize that knowledge to novel instances of those concepts, being able to flexibly transfer one's knowledge to new situations, robust reasoning using abstractions, analogies, and counterfactuals, and the ability to explain one's reasoning, among other abilities [14, 20, 39, 43].

Indeed, these are precisely the abilities lacking in current AI systems, including state-of-the-art LLMs, though ever-larger LLMs have exhibited limited sparks of these general abilities. The key questions of the debate about understanding in LLMs are the following: (1) Is talking of understanding in such systems simply a category error, namely, that these models are not, and will never be, the kind of things that can understand? Or conversely, (2) do these systems (or will their near-term successors) actually create something like the compressed "theories" that are central to human understanding, and, if so, does scaling these models create ever better theories? Or (3) If these systems do not create such compressed theories, can their unimaginably large systems of statistical correlations produce abilities that are functionally equivalent to human understanding, that is "competence without comprehension" [15]? And at this point will it still make sense to call such correlations "spurious" or the resulting solutions "shortcuts?"

While adherents on both sides of the "LLM understanding" debate have strong intuitions supporting their views, our current sciences of intelligence are inadequate for answering such questions. And these questions are no longer in the realm of abstract philosophical discussions, but touch on very real concerns about the capabilities, robustness, safety, and ethics of AI systems that are increasingly playing roles in humans' everyday lives.

The debate over understanding in LLMs, as ever larger and seemingly more capable systems are developed, underscores the need for a new science of intelligence that makes sense of broader conceptions of understanding and that can answer questions like those posed above, for both humans and machines. As neuroscientist Terrence Sejnowski points out, "The diverging opinions of experts on the intelligence of LLMs suggests that our old ideas based on natural intelligence are inadequate" [41]. If LLMs and related models succeed by exploiting statistical correlations at a heretofore unthinkable scale, perhaps this could be considered a novel form of "understanding", one that enables extraordinary, superhuman predictive ability, such as in the case of the AlphaZero and AlphaFold systems from DeepMind [27, 42], which respectively seem to bring an "alien" form of intuition to the domains of chess playing and protein structure prediction [26, 40].

It could thus be argued that in recent years the field of AI has created machines with new modes of understanding, most likely new species in a larger zoo of related concepts, that will continue to be enriched as we make progress in our pursuit of the elusive nature of intelligence. And just as different species are better adapted to different environments, our intelligent systems will be

5

better adapted to different problems. Problems that require enormous quantities of historically encoded knowledge where performance is at a premium will continue to favor large-scale statistical models like LLMs, and those for which we have limited knowledge and strong causal mechanisms will favor human intelligence. The challenge for the future is to develop new scientific methods that can reveal the detailed mechanisms of understanding in distinct forms of intelligence, discern their strengths and limitations, and learn how to integrate such truly diverse modes of cognition.

Acknowledgments

References

[1] B. Aguera y Arcas. Do large language models understand us?, 2021. Medium, December 16 `tinyurl.com/38t23n73`.

[2] B. Aguera y Arcas. Can machines learn how to behave?, 2022. Medium, August 3, `tinyurl.com/mr4cb3dw`.

[3] B. Aguera y Arcas. Artificial neural networks are making strides towards consciousness, 2022. The Economist, June 13, `tinyurl.com/ymhk37uu`.

[4] C. Baumberger, C. Beisbart, and G. Brun. What is understanding? an overview of recent debates in epistemology and philosophy of science. In *Explaining Understanding: New Perspectives from Epistemology and Philosophy of Science*, pages 1–34. Routledge, 2017.

[5] E. M. Bender and A. Koller. Climbing towards NLU: On meaning, form, and understanding in the age of data. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5185–5198, 2020.

[6] E. M. Bender, T. Gebru, A. McMillan-Major, and S. Shmitchell. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pages 610–623, 2021.

[7] R. Bommasani, D. A. Hudson, E. Adeli, R. Altman, S. Arora, S. von Arx, M. S. Bernstein, J. Bohg, A. Bosselut, E. Brunskill, et al. On the opportunities and risks of foundation models, 2021. arXiv:2108.07258.

[8] T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, et al. Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33:1877–1901, 2020.

[9] J. Browning and Y. LeCun. AI and the limits of language, 2022. Noema, August 23, `https://www.noemamag.com/ai-and-the-limits-of-language/`.

[10] S. R. Choudhury, A. Rogers, and I. Augenstein. Machine reading, fast and slow: When do models 'understand' language?, 2022. arXiv:2209.07430.

[11] A. Chowdhery, S. Narang, J. Devlin, M. Bosma, G. Mishra, A. Roberts, P. Barham, H. W. Chung, C. Sutton, S. Gehrmann, et al. PaLM: Scaling language modeling with pathways, 2022. arXiv:2204.02311.

[12] I. Dasgupta, A. K. Lampinen, S. C. Y. Chan, A. Creswell, D. Kumaran, J. L. McClelland, and F. Hill. Language models show human-like content effects on reasoning, 2022. arXiv:2207.07051.

[13] N. de Freitas, 2022. May 14, `https://twitter.com/NandoDF/status/1525397036325019649`.

[14] H. W. De Regt. Discussion note: Making sense of understanding. *Philosophy of Science*, 71 (1):98–109, 2004.

[15] D. C. Dennett. 'A perfect and beautiful machine': What Darwin's theory of evolution reveals about artificial intelligence, 2012. The Atlantic, June 22.

[16] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, page 4171–4186, 2019.

[17] A. Dimakis, 2022. May 16, `https://twitter.com/AlexGDimakis/status/1526388274348150784`.

[18] M. Gardner, W. Merrill, J. Dodge, M. E. Peters, A. Ross, S. Singh, and N. Smith. Competency problems: On finding and removing artifacts in language data. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, 2021.

[19] R. Geirhos, J.-H. Jacobsen, C. Michaelis, R. Zemel, W. Brendel, M. Bethge, and F. A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020.

[20] D. George, M. Lázaro-Gredilla, and J. S. Guntupalli. From CAPTCHA to commonsense: How brain can teach us about artificial intelligence. *Frontiers in Computational Neuroscience*, 14: 554097, 2020.

[21] A. Gopnik. What AI still doesn't know how to do, 2022. Washington Post, July 15, https://www.wsj.com/articles/what-ai-still-doesnt-know-how-to-do-11657891316.

[22] A. Gopnik and H. M. Wellman. The theory theory. In *Domain Specificity in Cognition and Culture*, pages 257–293. 1994.

[23] S. Gururangan, S. Swayamdipta, O. Levy, R. Schwartz, S. R. Bowman, and N. A. Smith. Annotation artifacts in natural language inference data. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 107–112, 2018.

[24] I. Habernal, H. Wachsmuth, I. Gurevych, and B. Stein. The argument reasoning comprehension task: Identification and reconstruction of implicit warrants. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, page 1930–1940, 2018.

[25] D. R. Hofstadter. *Fluid Concepts and Creative Analogies: Computer Models of the Fundamental Mechanisms of Thought*. Basic Books, 1995. Preface to Chapter 4.

[26] D. T. Jones and J. M. Thornton. The impact of AlphaFold2 one year on. *Nature Methods*, 19 (1):15–20, 2022.

[27] J. Jumper, R. Evans, A. Pritzel, T. Green, M. Figurnov, O. Ronneberger, K. Tunyasuvunakool, R. Bates, A. Žídek, A. Potapenko, et al. Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873):583–589, 2021.

[28] J. L. Kvanvig. Knowledge, understanding, and reasons for belief. In *The Oxford Handbook of Reasons and Normativity*, page 685–705. Oxford University Press, 2018.

[29] S. Lapuschkin, S. Wäldchen, A. Binder, G. Montavon, W. Samek, and K.-R. Müller. Unmasking Clever Hans predictors and assessing what machines really learn. *Nature Communications*, 10(1):1–8, 2019.

[30] A. Laverghetta, A. Nighojkar, J. Mirzakhalov, and J. Licato. Predicting human psychometric properties using computational language models. In *Annual Meeting of the Psychometric Society*, pages 151–169. Springer, 2022.

[31] T. Linzen. How can we accelerate progress towards human-like linguistic generalization? In *In Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, page 5210–17, 2020.

[32] C. D. Manning. Human language understanding and reasoning. *Daedalus*, 151(2):127–138, 2022.

[33] G. Marcus. Nonsense on stilts, 2022. Substack, June 12, https://garymarcus.substack.com/p/nonsense-on-stilts.

[34] R. T. McCoy, E. Pavlick, and T. Linzen. Right for the wrong reasons: Diagnosing syntactic heuristics in natural language inference. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, page 3428–3448, 2019.

[35] J. Michael, A. Holtzman, A. Parrish, A. Mueller, A. Wang, A. Chen, D. Madaan, N. Nangia, R. Y. Pang, J. Phang, et al. What do NLP researchers believe? Results of the NLP community metasurvey, 2022. arXiv:2208.12852.

[36] B. Min, H. Ross, E. Sulem, A. P. B. Veyseh, T. H. Nguyen, O. Sainz, E. Agirre, I. Heinz, and D. Roth. Recent advances in natural language processing via large pre-trained language models: A survey, 2021. arXiv:2111.01243.

[37] M. Mitchell. Artificial intelligence hits the barrier of meaning. *Information*, 10(2):51, 2019.

[38] T. Niven and H.-Y. Kao. Probing neural network comprehension of natural language arguments. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4658–4664. Association for Computational Linguistics, 2019.

[39] J. Pearl. Theoretical impediments to machine learning with seven sparks from the causal revolution, 2018. arXiv:1801.04016.

[40] M. Sadler and N. Regan. *Game changer: AlphaZero's Groundbreaking Chess Strategies and the Promise of AI*. Alkmaar, 2019.

[41] T. Sejnowski. Large language models and the reverse Turing test, 2022. arXiv:2207.14382.

[42] D. Silver, T. Hubert, J. Schrittwieser, I. Antonoglou, M. Lai, A. Guez, M. Lanctot, L. Sifre, D. Kumaran, T. Graepel, et al. Mastering chess and shogi by self-play with a general reinforcement learning algorithm, 2017. arXiv:1712.01815.

[43] M. Strevens. No understanding without explanation. *Studies in History and Philosophy of Science Part A*, 44(3):510–515, 2013.

[44] R. Thoppilan, D. De Freitas, J. Hall, N. Shazeer, A. Kulshreshtha, H.-T. Cheng, A. Jin, T. Bos, L. Baker, Y. Du, et al. LaMDA: Language models for dialog applications, 2022. arXiv:2201.08239.

[45] S. Trott, C. Jones, T. Chang, J. Michaelov, and B. Bergen. Do large language models know what humans know?, 2022. arxiv:2209.01515.

[46] A. Wang, A. Singh, J. Michael, F. Hill, O. Levy, and S. R. Bowman. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 353–355. Association for Computational Linguistics, 2018.

[47] A. Wang, Y. Pruksachatkun, N. Nangia, A. Singh, J. Michael, F. Hill, O. Levy, and S. R. Bowman. SuperGLUE: A stickier benchmark for general-purpose language understanding systems. In *Advances in Neural Information Processing Systems*, volume 32, pages 3266–3280, 2019.

[48] J. Wei, Y. Tay, R. Bommasani, C. Raffel, B. Zoph, S. Borgeaud, D. Yogatama, M. Bosma, D. Zhou, D. Metzler, et al. Emergent abilities of large language models, 2022. arXiv:2206.07682.

[49] J. Wei, X. Wang, D. Schuurmans, M. Bosma, E. Chi, Q. Le, and D. Zhou. Chain of thought prompting elicits reasoning in large language models, 2022. arXiv:2201.11903.

[50] J. Weizenbaum. *Computer Power and Human Reason: From Judgment to Calculation.* WH Freeman & Co, 1976.